

Score - Based Generative Modelling

Manos Plitsis

What is Generative Modelling?

Given a dataset $\{x_i \in \mathbb{R}^D\}_{i=1}^N$, model the *data distribution* $p_{data}(x)$.

Once we have $p(x)$ we can generate new data points by sampling from it.

Energy-Based Modelling

To model an arbitrarily flexible distribution $p(x)$, we can model it as:

$$p_{\theta}(x) = \frac{e^{-f_{\theta}(x)}}{Z_{\theta}}$$

e.g. with maximum likelihood.

$f_{\theta}(x)$ is called the energy function

$Z_{\theta} = \int e^{-f_{\theta}(x)}$ is a normalizing constant

Energy-Based Modelling

To model an arbitrarily flexible distribution $p(x)$, we can model it as:

$$p_{\theta}(x) = \frac{e^{-f_{\theta}(x)}}{Z_{\theta}}$$

e.g. with maximum likelihood.

But Z is not tractable!



What if we could get rid of Z?

$$\begin{aligned}\nabla_x \log p_\theta(x) &= \nabla_x \log\left(\frac{e^{-f_\theta(x)}}{Z_\theta}\right) \\ &= \nabla_x \log\left(\frac{1}{Z_\theta}\right) + \nabla_x \log(e^{-f_\theta(x)}) \\ &= -\nabla_x f_\theta(x) \\ &\approx s_\theta(x)\end{aligned}$$

Score Networks

Def. The (*Stein*) score of a probability density $p(x)$ is $\nabla_x \log p(x)$

A score network $s_\theta(x) : \mathbb{R}^D \rightarrow \mathbb{R}^D$ is a neural network trained to approximate the score of $p_{data}(x)$




Score Networks

Def. The (Stein) score of a probability density $p(x)$ is $s(x) = \nabla_x \log p(x)$

A score network $s_\theta(x) : \mathbb{R}^D \rightarrow \mathbb{R}^D$ is a neural network trained to approximate the score of $p_{data}(x)$

Q1: How do we learn $s_\theta(x)$?

Q2: How do we use it to generate new samples?



Score Matching

Q1. How do we learn $s_\theta(x)$?

A. We optimize the score matching objective:

$$L(\theta) \triangleq \frac{1}{2} \mathbb{E}_{p_{data}} [\|s_\theta(x) - \nabla_x \log p_{data}(x)\|_2^2]$$

Score Matching

Q1: How do we learn $s_\theta(x)$?

A: We optimize the score matching objective:

$$L(\theta) \triangleq \frac{1}{2} \mathbb{E}_{p_{data}} [\|s_\theta(x) - \nabla_x \log p_{data}(x)\|_2^2]$$

But we don't have access to $\nabla_x \log p_{data}(x)$!

L can be shown equivalent (up to a constant) to:

$$\mathbb{E}_{p_{data}(x)} [tr(\nabla_x s_\theta(x)) + \frac{1}{2} \|s_\theta(x)\|_2^2]$$

We can compute this from data, but calculating $tr(\nabla_x s_\theta(x))$ is too costly.

Score Matching

Denoising Score Matching

We perturb the data with a noise distribution $q_\sigma(\tilde{x}|x)$, then estimate the score of $q_\sigma \triangleq \int q_\sigma(\tilde{x}|x)p_{data}(x)dx$ with objective:

$$\frac{1}{2} \mathbb{E}_{q_\sigma(\tilde{x}|x)p_{data}(x)} [\|s_\theta(\tilde{x}) - \nabla_{\tilde{x}} \log q_\sigma(\tilde{x}|x)\|_2^2]$$

E.g. for Gaussian $\nabla_{\tilde{x}} \log q_\sigma(\tilde{x}|x) = -\frac{\tilde{x}-x}{\sigma^2}$

Then $s_{\theta_*}(x) = \nabla_x \log q_\sigma(x)$ almost surely.

But: $s_{\theta_*}(x) = \nabla_x \log q_\sigma(x) \approx \nabla_x \log p_{data}(x)$ only when σ is small s.t. $q_\sigma(x) \approx p_{data}(x)$ (But not too small and the variance explodes)

Langevin Dynamics

Q2: How do we use $s_\theta(x)$ to generate new samples?

A: We recursively compute:



$$\tilde{x}_t = \tilde{x}_{t-1} + \frac{\epsilon}{2} \nabla_x \log p(\tilde{x}_{t-1}) + \sqrt{\epsilon} z_t$$

where

- $\epsilon > 0$ a fixed step size
- $\tilde{x}_0 \sim \pi(x)$ with π a prior distribution
- $z_t \sim \mathcal{N}(0, I)$

The distribution of \tilde{x}_T approaches $p(x)$ when $\epsilon \rightarrow 0, T \rightarrow \infty$

Langevin Dynamics



Challenges of Score-Based Modelling

1. Data are concentrated on a low dimensional manifold
2. Scarcity of data leads to low density regions

1. The manifold hypothesis

Real world data tend to concentrate on **low dimensional manifolds** embedded in the ambient space.

- Score is undefined outside the manifold.
- Score matching only works when the support of $p(x)$ is the ambient space
[Hyvärinen 2005]


2. Scarcity of data

Inaccurate score estimation


In practice, $\mathbb{E}_{p_{data}}$ is estimated using i.i.d. samples.

For regions $\mathcal{R} \subset \mathbb{R}^D$ s.t $p_{data}(\mathcal{R}) \approx 0$ there usually isn't enough data to estimate the score.

Data scores



Estimated scores



2. Scarcity of data


Slow mixing of Langevin Dynamics

Let $p_1(x), p_2(x)$ normalized distributions with disjoint support, and
 $p_{data}(x) = \pi p_1(x) + (1 - \pi)p_2(x)$, $\pi \in (0, 1)$


The score does not depend on π , so samples with LD will not depend on π .

In practice, also true when supports are *approximately disjoint*.

i.i.d samples



Langevin dynamics samples



Noise Conditional Score Networks

with Annealed Langevin Dynamics

Idea:

- Perturb data with various noise levels $\sigma_1, \dots, \sigma_L$ using distribution

$$q_\sigma(x) \triangleq \int p_{data}(t) \mathcal{N}(x|t, \sigma^2 I) dt$$

(this gives $q_\sigma(\tilde{x}|x) = \mathcal{N}(\tilde{x}|x, \sigma^2 I)$, the "noising" distribution)

- Train a single conditional score network $s_\theta(x, \sigma)$ for all noise levels σ_i to predict the noised data score $\nabla x \log q_{\sigma_i}(x)$.
- In LD, start with high noise level scores, and gradually decrease the noise level.

Noise Conditional Score Networks

- $\frac{\sigma_1}{\sigma_2} = \dots = \frac{\sigma_{L-1}}{\sigma_L} > 1$
- σ_1 is large enough to avoid previous problems.
- σ_L is small enough to minimize effect on data.

Noise Conditional Score Networks

We defined $q_\sigma(\tilde{x}|x) = \mathcal{N}(\tilde{x}|x, \sigma^2 I)$,
so $\nabla_{\tilde{x}} \log q_\sigma(\tilde{x}|x) = -\frac{\tilde{x}-x}{\sigma^2}$.

For a given σ :

$$\ell(\theta; \sigma) \triangleq \frac{1}{2} \mathbb{E}_{p_{\text{data}}(x)} \mathbb{E}_{\tilde{x} \sim \mathcal{N}(x, \sigma^2 I)} \left[\left\| s_\theta(\tilde{x}, \sigma) + \frac{\tilde{x} - x}{\sigma^2} \right\|_2^2 \right]$$

So the final loss is:

$$\mathcal{L}(\theta; \{\sigma_i\}_{i=1}^L) \triangleq \frac{1}{L} \sum_{i=1}^L \lambda(\sigma_i) \ell(\theta; \sigma_i)$$

Noise Conditional Score Networks

How to choose $\lambda(\sigma)$?

- We want $\lambda(\sigma_i)\ell(\theta; \sigma_i)$ to have the same order of magnitude for all $\{\sigma_i\}_{i=1}^L$.
- With optimally trained networks, we observe $\|s_\theta(x, \sigma)\|_2 \propto \frac{1}{\sigma}$.
- This inspires the choice $\lambda(\sigma) = \sigma^2$, leading to
$$\lambda(\sigma)\ell(\theta; \sigma) = \sigma^2\ell(\theta; \sigma) = \frac{1}{2} \left\| \sigma s_\theta(x, \sigma) + \frac{\hat{x} - x}{\sigma} \right\|_2^2.$$
- Given $\frac{\hat{x} - x}{\sigma} \sim \mathcal{N}(0, I)$ and $\|\sigma s_\theta(x, \sigma)\|_2 \propto 1$, the magnitude of $\lambda(\sigma)\ell(\theta; \sigma)$ does not depend on σ .


Annealed Langevin Dynamics

Algorithm 1: Annealed Langevin dynamics

Require : $\{\sigma_i\}_{i=1}^L, \epsilon, T$.

- 1 : Initialize \tilde{x}_0
- 2 : **for** $i = 1$ **to** L **do**
- 3 : $\alpha_i \leftarrow \epsilon \sigma_i^2 / \sigma_L^2$ (α_i is the step size)
- 4 : **for** $t = 1$ **to** T **do**
- 5 : Draw $z_t \sim \mathcal{N}(0, I)$
- 6 : $\tilde{x}_t \leftarrow \tilde{x}_{t-1} + \frac{\alpha_i}{2} s_\theta(\tilde{x}_{t-1}, \sigma_i) + \sqrt{\alpha_i} z_t$
- 7 : **end for**
- 8 : $\tilde{x}_0 \leftarrow \tilde{x}_T$
- 9 : **end for**
- 10 : **return** \tilde{x}_T

Annealed Langevin Dynamics




Annealed Langevin Dynamics


How to choose α_i ?

- Many ways to tune α_i . We use $\alpha_i \propto \sigma_i^2$.
- Aim: Fix the magnitude of the "signal-to-noise" ratio $\frac{\alpha_i s_\theta(x, \sigma_i)}{2\sqrt{\alpha_i}z}$ in Langevin dynamics w.r.t. σ_i .
- $\mathbb{E} \left[\left\| \frac{\alpha_i s_\theta(x, \sigma_i)}{2\sqrt{\alpha_i}z} \right\|_2^2 \right] \approx \mathbb{E} \left[\frac{\alpha_i \|s_\theta(x, \sigma_i)\|_2^2}{4} \right] \propto \frac{1}{4} \mathbb{E} \left[\|\sigma_i s_\theta(x, \sigma_i)\|_2^2 \right]$.
- Empirically, when networks are optimally trained, $\|s_\theta(x, \sigma_i)\|_2 \propto \frac{1}{\sigma_i}$.
- Thus $\mathbb{E} \left[\|\sigma_i s_\theta(x, \sigma_i)\|_2^2 \right] \propto 1$.
- Therefore, the choice of α_i does not depend on σ_i .

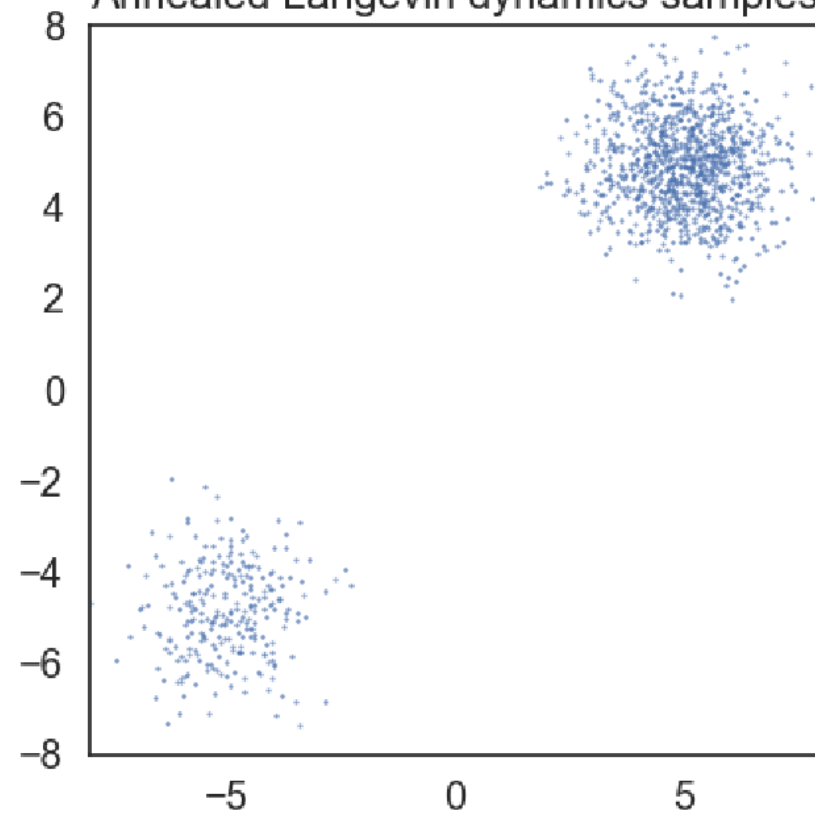
i.i.d samples



Langevin dynamics samples



Annealed Langevin dynamics samples




Experimental Setup

- $L = 10$
- $\{\sigma_i\}_{i=1}^L$ is a geometric sequence with:
 - $\sigma_1 = \sigma_{max} = 1$
 - $\sigma_{10} = \sigma_{min} = 0.01$
- For LD sampling:
 - $T = 10$
 - $\epsilon = 2 \times 10^{-5}$
 - Initial samples are uniform noise.

Results

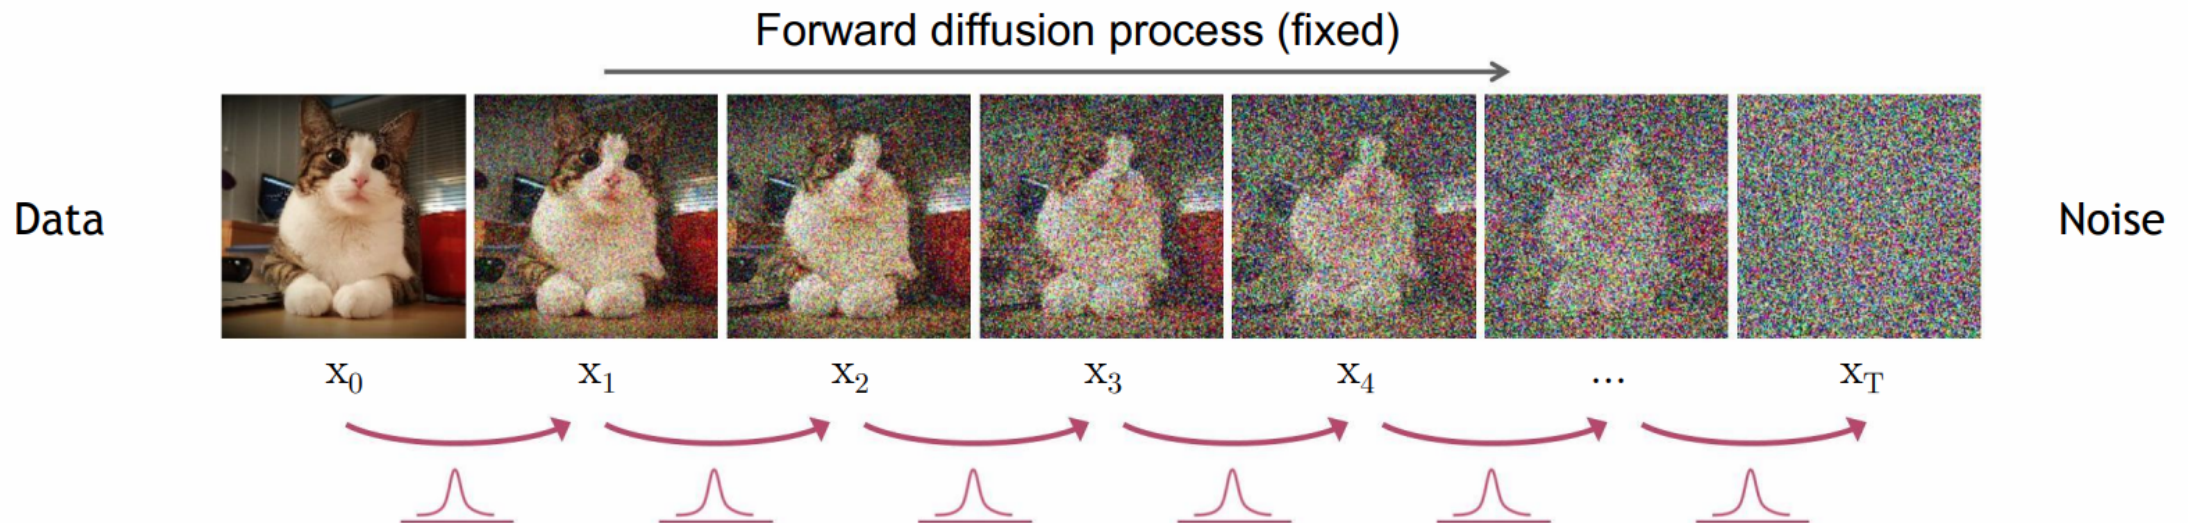
Model	Inception	FID
CIFAR-10 Unconditional		
PixelCNN [59]	4.60	65.93
PixelIQN [42]	5.29	49.46
EBM [12]	6.02	40.58
WGAN-GP [18]	$7.86 \pm .07$	36.4
MoLM [45]	$7.90 \pm .10$	18.9
SNGAN [36]	$8.22 \pm .05$	21.7
ProgressiveGAN [25]	$8.80 \pm .05$	-
NCSN (Ours)	$8.87 \pm .12$	25.32
CIFAR-10 Conditional		
EBM [12]	8.30	37.9
SNGAN [36]	$8.60 \pm .08$	25.5
BigGAN [6]	9.22	14.73



NCSM vs DDPM

Reminder: DDPM

- $q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{\alpha_t}x_{t-1}, (1 - \alpha_t)I)$ (noising step)
- $p_\theta(x_T) = \mathcal{N}(x_T; 0, I)$.
- $p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \sigma(x_t, t)^2 I)$ (denoising step)



ELBO for DDPM

$$\begin{aligned}\log p(x) &\geq \mathbb{E}_{q(x_1|x_0)} [\log p_\theta(x_0|x_1)] && (L_0 : \text{Reconstruction term}) \\ &\quad - D_{KL}(q(x_T|x_0)||p(x_T)) && (L_T : \text{Prior matching term}) \\ &\quad - \sum_{t=2}^T \mathbb{E}_{q(x_t|x_0)} [D_{KL}(q(x_{t-1}|x_t, x_0)||p_\theta(x_{t-1}|x_t))] && (L_{t-1} : \text{Denoising matching term})\end{aligned}$$

To **maximize** the ELBO, we need to minimize the denoising matching term.

We can write $x_t = \sqrt{\bar{a}_t}x_0 + \sqrt{1 - \bar{a}_t}\epsilon_0 \sim \mathcal{N}(x_t; \sqrt{\bar{a}_t}x_0, (1 - \bar{a}_t)I)$

with $\bar{a}_t = \prod_{i=1}^t a_i$

Then:

$$\begin{aligned}
 q(x_{t-1}|x_t, x_0) &= \frac{q(x_t|x_{t-1}, x_0)q(x_{t-1}|x_0)}{q(x_t|x_0)} \quad (\text{Bayes rule}) \\
 &= \frac{\mathcal{N}(x_t; \sqrt{\bar{a}_t}x_0, (1 - \bar{a}_t)I) \mathcal{N}(x_{t-1}; \sqrt{\bar{a}_{t-1}}x_0, (1 - \bar{a}_{t-1})I)}{\mathcal{N}(x_t; \sqrt{\bar{a}_t}x_0, (1 - \bar{a}_t)I)} \\
 &= \dots \\
 &= \mathcal{N}\left(x_{t-1}; \underbrace{\frac{\sqrt{\bar{a}_t}(1 - \bar{a}_{t-1})x_t + \sqrt{\bar{a}_{t-1}}(1 - a_t)x_0}{1 - \bar{a}_t}}_{\mu_q(x_t, x_0)}, \underbrace{\frac{(1 - a_t)(1 - \bar{a}_{t-1})}{(1 - \bar{a}_t)}I}_{\Sigma_q(t)}\right)
 \end{aligned}$$

Learning μ_θ

$$\begin{aligned} & \operatorname{argmin}_\theta D_{KL}(q(x_{t-1}|x_t, x_0) \parallel p_\theta(x_{t-1}|x_t)) \\ = & \operatorname{argmin}_\theta D_{KL}(\mathcal{N}(x_{t-1}; \mu_q(t), \Sigma_q(t)) \parallel \mathcal{N}(x_{t-1}; \mu_\theta(t), \Sigma_q(t))) && \text{(set denoising transition variance to be } \Sigma_q(t)) \\ = & \dots && \text{(KL Divergence Gaussians)} \\ = & \operatorname{argmin}_\theta \frac{1}{2\sigma_q^2(t)} [\|\mu_\theta - \mu_q\|_2^2] \end{aligned}$$

Learning ϵ_θ

We can choose the parameterization: $x_0 = \frac{x_t + \sqrt{1 - \bar{a}_t} \epsilon_0}{\sqrt{\bar{a}_t}}$

- $\mu_q(x_t, x_0) = \frac{1}{\sqrt{a_t}} x_t - \frac{1 - a_t}{\sqrt{1 - \bar{a}_t} \sqrt{a_t}} \epsilon_0$
- $\mu_\theta(x_t, t) = \frac{1}{\sqrt{a_t}} x_t - \frac{1 - a_t}{\sqrt{1 - \bar{a}_t} \sqrt{a_t}} \epsilon_\theta(x_t, t)$

Reformulate the loss to:

$$\operatorname{argmin}_\theta = \underbrace{\frac{(1 - a_t)^2}{2\sigma_q^2(t)(1 - \bar{a}_t)a_t}}_{\lambda_t} [\|e_0 - e_\theta(x_t, t)\|_2^2]$$

Learning $\nabla_x \log p_\theta(x)$

Given a Gaussian variable $z \sim \mathcal{N}(z; \mu_z, \Sigma_z)$, Tweedie's Formula states:

$$\mathbb{E}[\mu_z|z] = z + \Sigma_z \nabla_z \log p(z)$$

From a known equation, we have:

$$q(x_t|x_0) = \mathcal{N}(x_t; \sqrt{\alpha_t}x_0, (1 - \alpha_t)I)$$

By Tweedie's Formula, we get:

$$\mathbb{E}[\mu_{x_t}|x_t] = x_t + (1 - \alpha_t) \nabla_{x_t} \log p(x_t)$$

The best estimate for the true mean $\mu_{x_t} = \sqrt{\alpha_t}x_0$, is:

$$\begin{aligned} \sqrt{\alpha_t}x_0 &= x_t + (1 - \alpha_t) \nabla_{x_t} \log p(x_t) \\ \Rightarrow x_0 &= x_t + \frac{(1 - \alpha_t)}{\sqrt{\alpha_t}} \nabla_{x_t} \log p(x_t) \end{aligned}$$

Learning $\nabla_x \log p_\theta(x)$

Remember,

$$\mu_q(x_t, x_0) = \frac{\sqrt{a_t}(1-\bar{a}_{t-1})x_t + \sqrt{\bar{a}_{t-1}}(1-a_t)\mathbf{x}_0}{1-\bar{a}_t} = \frac{\sqrt{a_t}(1-\bar{a}_{t-1})x_t + \sqrt{\bar{a}_{t-1}}(1-a_t)\left(x_t + \frac{(1-a_t)}{\sqrt{a_t}} \nabla_{x_t} \log p(x_t)\right)}{1-\bar{a}_t}$$

- $\mu_q(x_t, x_0) = \frac{1}{\sqrt{a_t}} x_t - \frac{1-a_t}{\sqrt{a_t}} \nabla_{x_t} \log p(x_t)$
- $\mu_\theta(x_t, t) = \frac{1}{\sqrt{a_t}} x_t - \frac{1-a_t}{\sqrt{a_t}} s_\theta(x_t, t)$

Reformulate the loss to:


$$\operatorname{argmin}_\theta = \frac{(1-a_t)^2}{2\sigma_q^2(t)a_t} \left[\|s_\theta(x_t, t) - \nabla_{x_t} \log p(x_t)\|_2^2 \right]$$

The score looks like ϵ_0 (!?)

$$x_0 = \frac{x_t + \sqrt{1 - \bar{\alpha}_t} \epsilon_0}{\sqrt{\bar{\alpha}_t}} = x_t + \frac{(1 - \alpha_t)}{\sqrt{\alpha_t}} \nabla_{x_t} \log p(x_t)$$
$$\Rightarrow \nabla_{x_t} \log p(x_t) = -\frac{1}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_0$$

Perturbing images with SDEs

When the noise levels approach infinity, we essentially perturb the data with a **Stochastic Differential Equation (SDE)**



SDEs

In general, SDEs have the form

$$dx = f(x, t)dt + g(t)dw$$

where $f(\cdot, t) : \mathbb{R}^D \rightarrow \mathbb{R}^D$ is called the *drift coefficient* and $g(t) \in \mathbb{R}$ is called the *diffusion coefficient*.


w denotes a standard Brownian motion, and dw can be considered an infinitesimal white noise.

The solution is a continuous collection of random variables $\{x(t)\}_{t \in [0,1]}$.

Reversing the SDE

To sample from $x(T) \sim p_T$ and get new data from p_{data} , we can reverse the SDE (the reverse of a diffusion process is also a diffusion process [Anderson 1982]):

$$dx = [f(x, t) - g^2(t) \nabla_x \log p_t(x)] dt + g(t) d\tilde{w}.$$



Estimating the scores

To solve the reverse-time SDE we need:

- The terminal distribution $p(T) \approx \pi(x)$
- The score $\nabla_x \log p_t(x)$

To estimate the score we can use score matching techniques to train a *time dependent score model*, with objective:

$$\mathbb{E}_{t \in U(0, T)} \mathbb{E}_{p_t(x)} [\lambda(t) \|\nabla_x \log p_t(x) - s_\theta(x, t)\|_2^2]$$

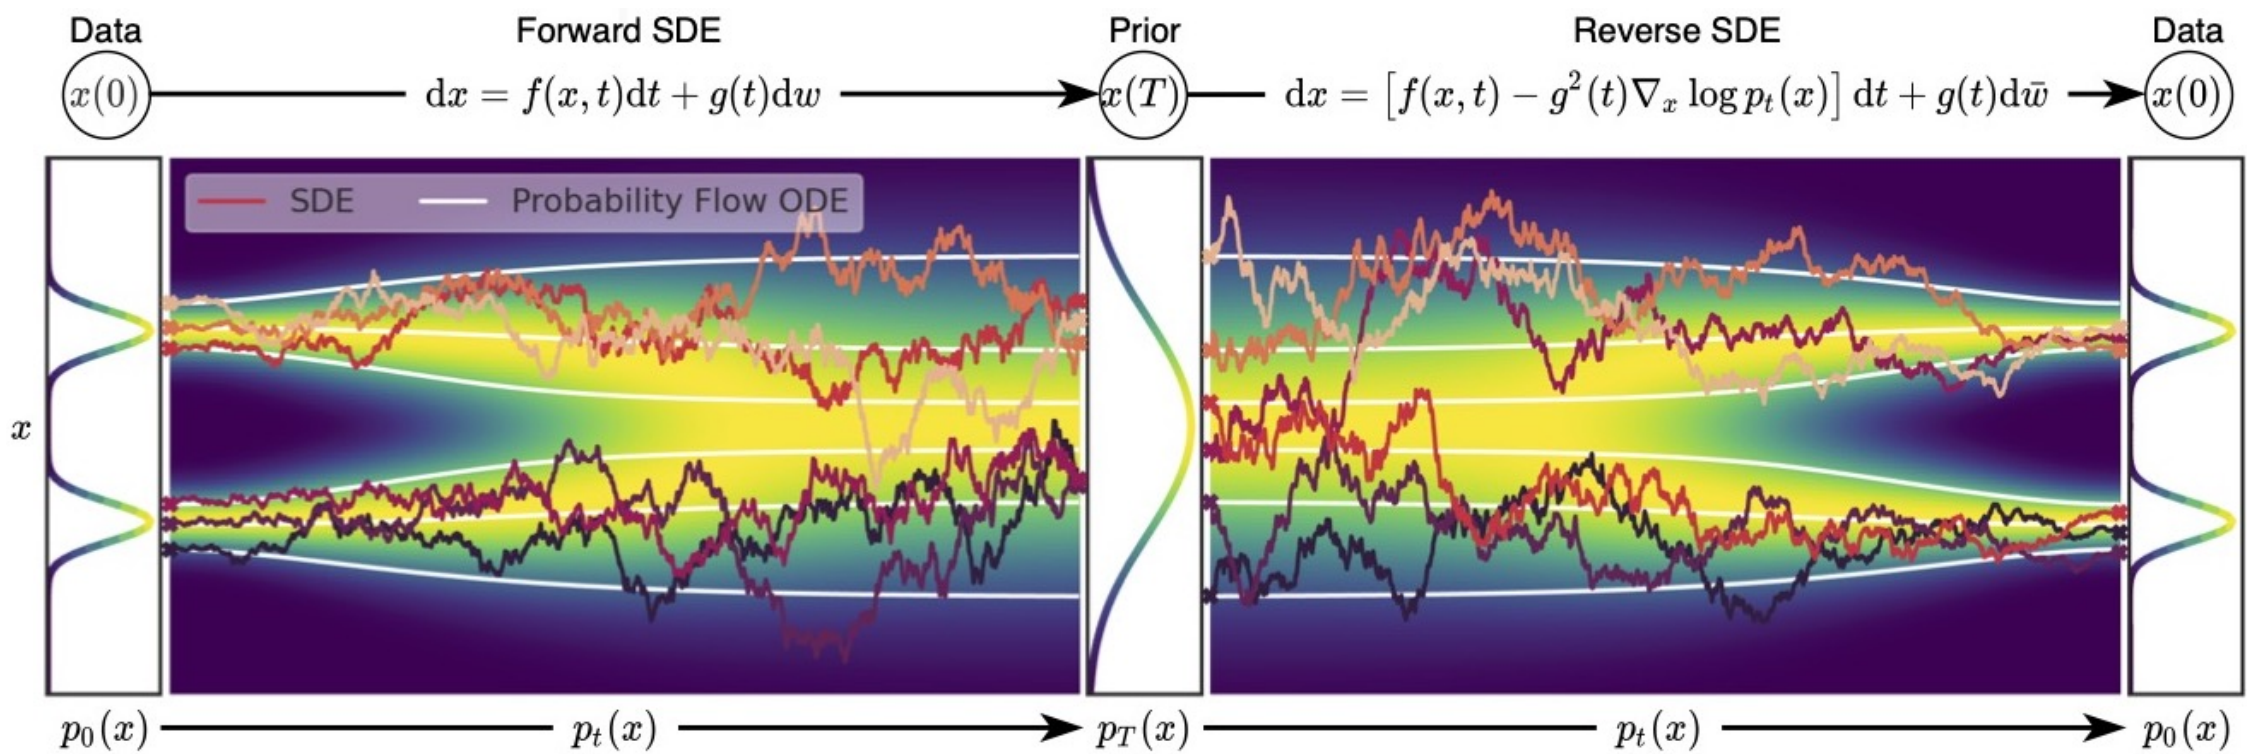
Then we can solve the reverse-time SDE with numerical SDE solvers (e.g. Euler - Maruyama)

Probability flow ODEs

For all diffusion processes, there exists a deterministic process whose trajectories share the same marginal probabilities $\{p_t(x)\}_{t=0}^T$ as the SDE.

This process satisfies the *probability flow ODE*:

$$dx = \left[f(x, t) - \frac{1}{2} g^2(t) \nabla_x \log p_t(x) \right] dt.$$



Neural ODE

When $\nabla_x \log p_t(x)$ is replaced by $s(x, t)$, it becomes a special case of *neural ODE*, specifically continuous normalizing flows.

So we get:

- Exact likelihood estimation.
- Encoding data points $x(0)$ to latent space $x(T)$.
 - Decoding by integrating corresponding ODE for reverse-time SDE.
 - We can manipulate the latent representation for editing by interpolation, temperature scaling.
- We get a uniquely identifiable encoding given sufficient data and model capacity.
- Efficient sampling by discretizing the ODE (link with DDIM?)

Thank you for your attention!

